

Pulp Tasking System 2025

Updates on the Pulp Project's kitchen

~ by Pedro Brochado

Foreword

- Thanks to the tasking working group!
 - Dennis
 - Brian
 - Decko
 - Matthias

Foreword

- Thanks to the tasking working group!
 - Dennis
 - Brian
 - Decko
 - Matthias
- Special thanks for Matthias for mentoring me in the tasking system and concurrency world

Overview

Overview

1. Background

Overview

1. Background
2. Cooking in Pulp's kitchen

Overview

1. Background
2. Cooking in Pulp's kitchen
3. New menu

Overview

1. Background
2. Cooking in Pulp's kitchen
3. New menu
4. Final words

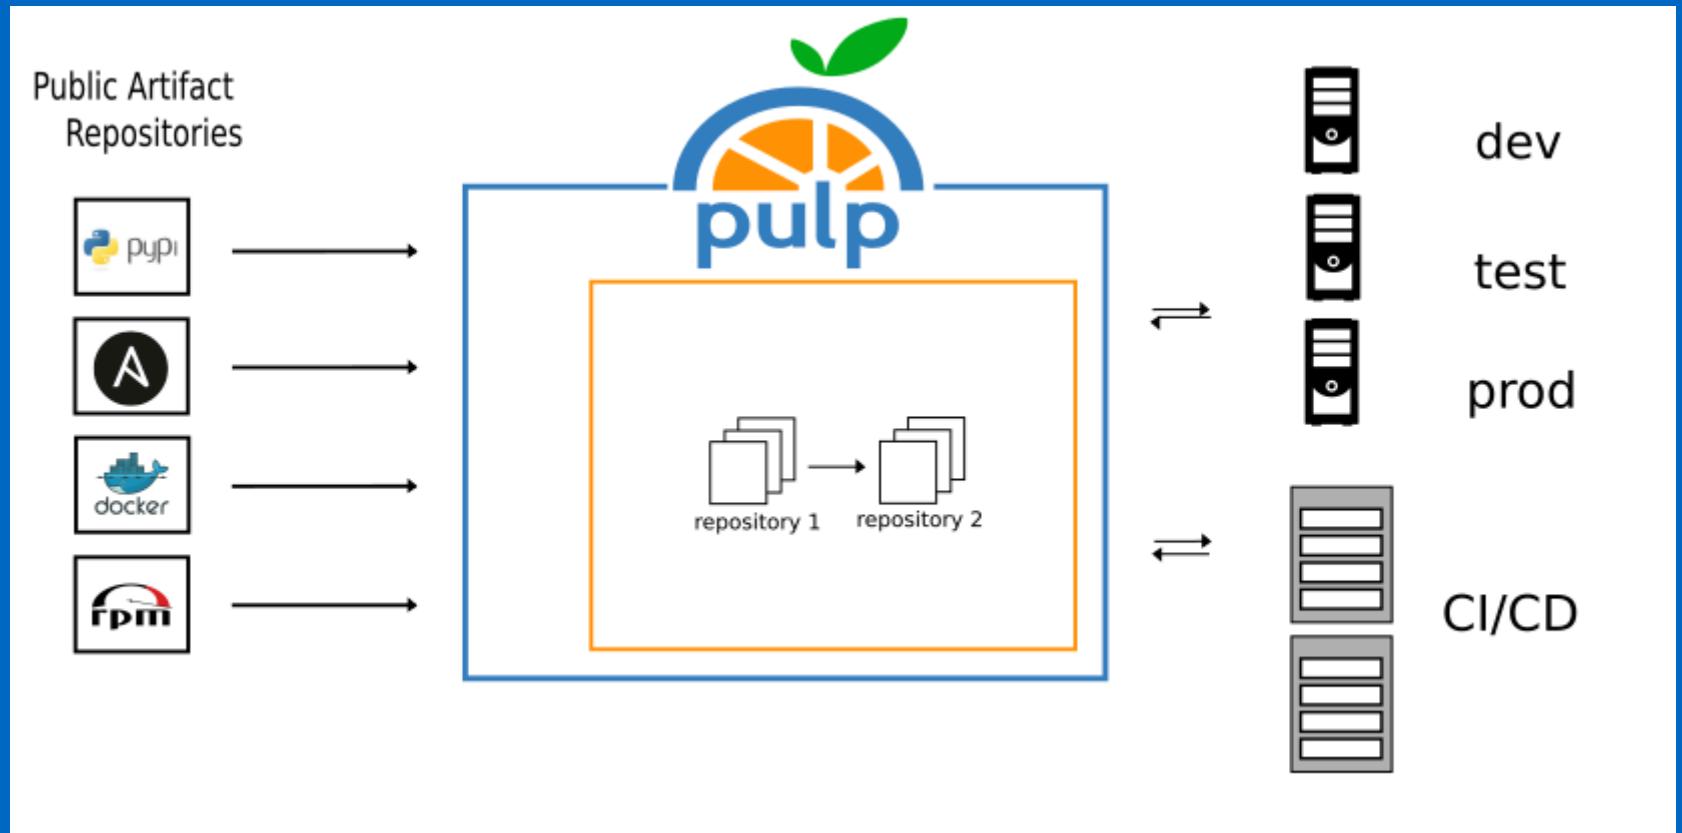
1. Background

Background • Pulp's Kitchen • New menu • Final words

What is Pulp?

Background • Pulp's Kitchen • New menu • Final words

What is Pulp?

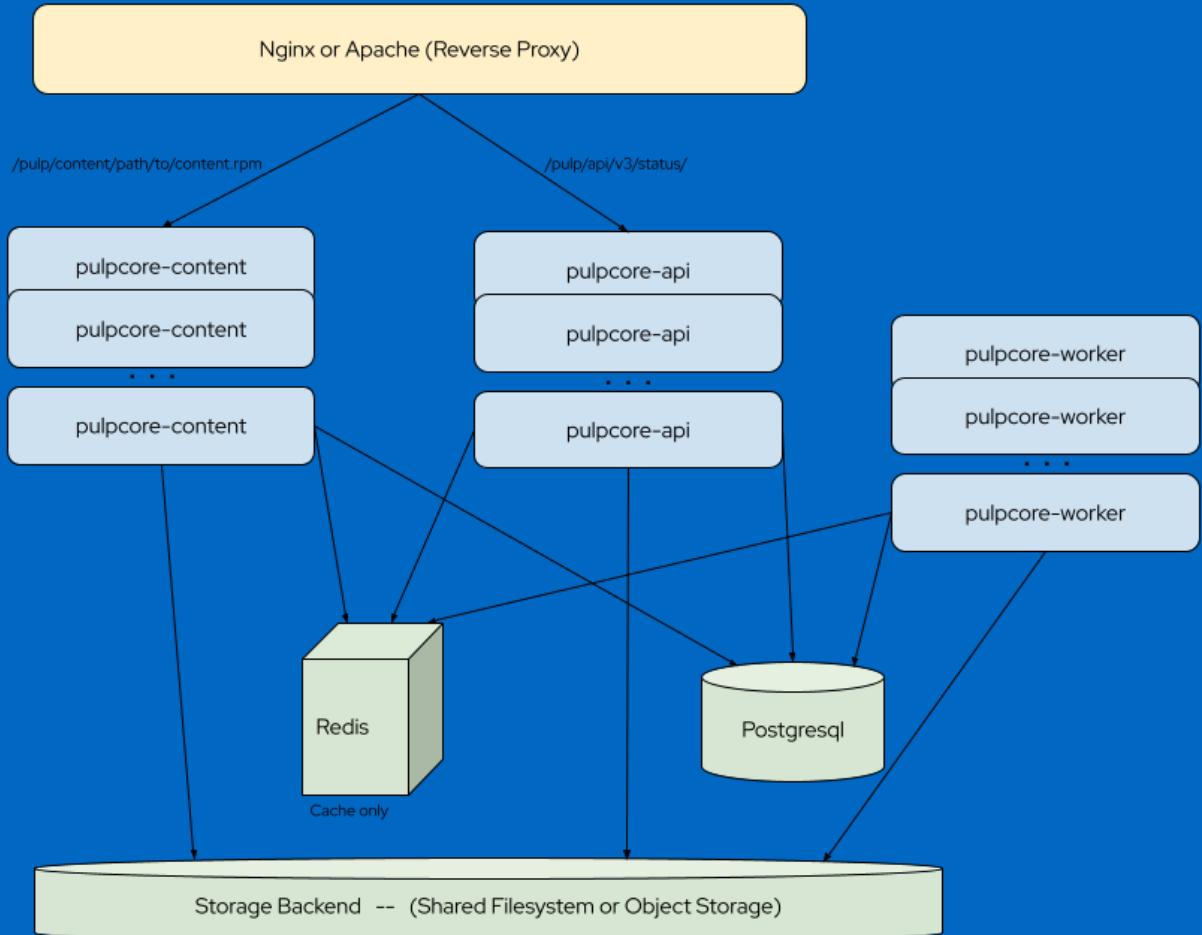


Background • Pulp's Kitchen • New menu • Final words

What is Pulp?

- An application for content management that can:
 - Serve content
 - Sync content from remote
 - Process uploaded content
- All these operations (and a lot more, of course) can trigger tasks!

Pulp's architecture



Background • Pulp's Kitchen • New menu • Final words

Why a tasking system?

Background • Pulp's Kitchen • New menu • Final words

Why a tasking system?

Handle long running tasks

Why a tasking system?

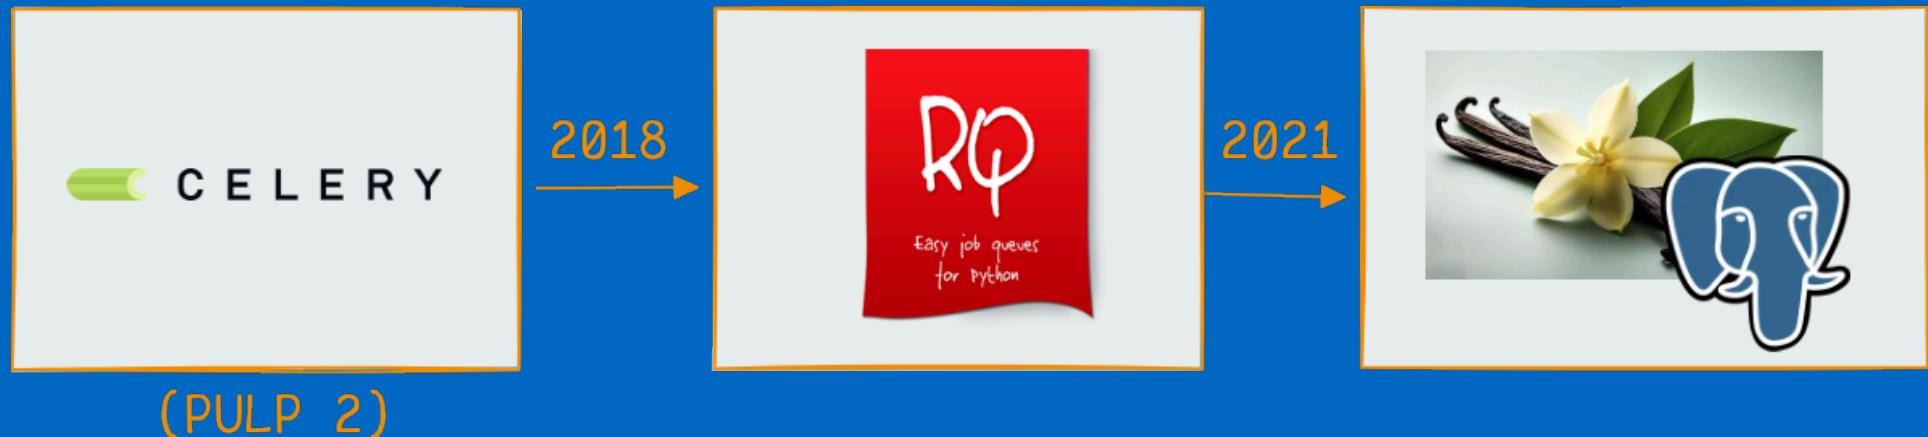
Handle long running tasks

Serialize operations where concurrency is unsafe

Previous systems

Background • Pulp's Kitchen • New menu • Final words

Previous systems

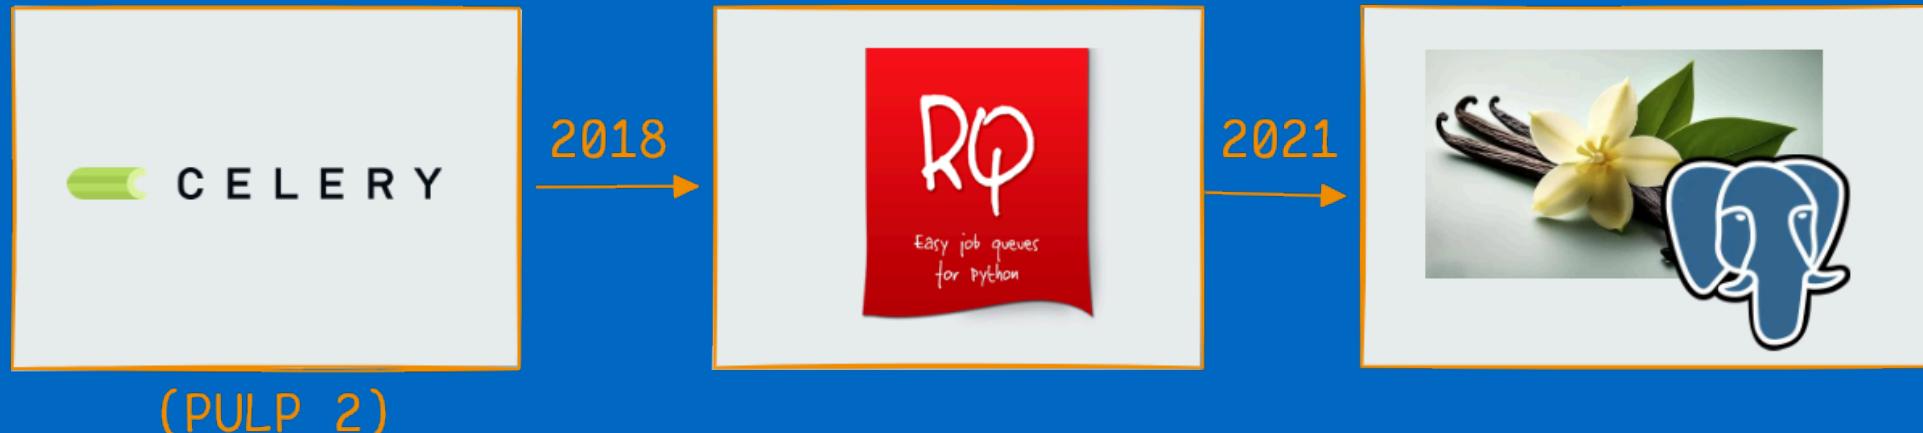


Background • Pulp's Kitchen • New menu • Final words

Previous systems

RQ pitfalls:

- Resource manager bottleneck (for task processing)
- Problems with data duplication (redis and postgres)



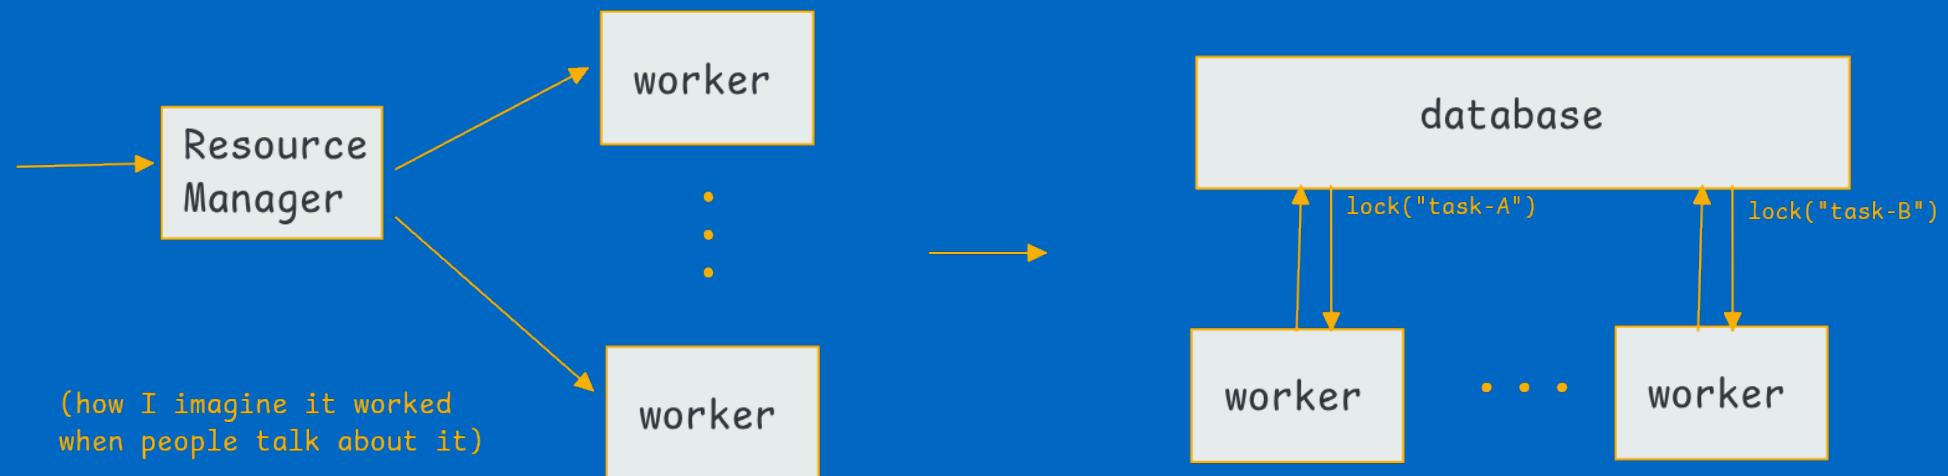
Background • Pulp's Kitchen • New menu • Final words

System in 2021

Background • Pulp's Kitchen • New menu • Final words

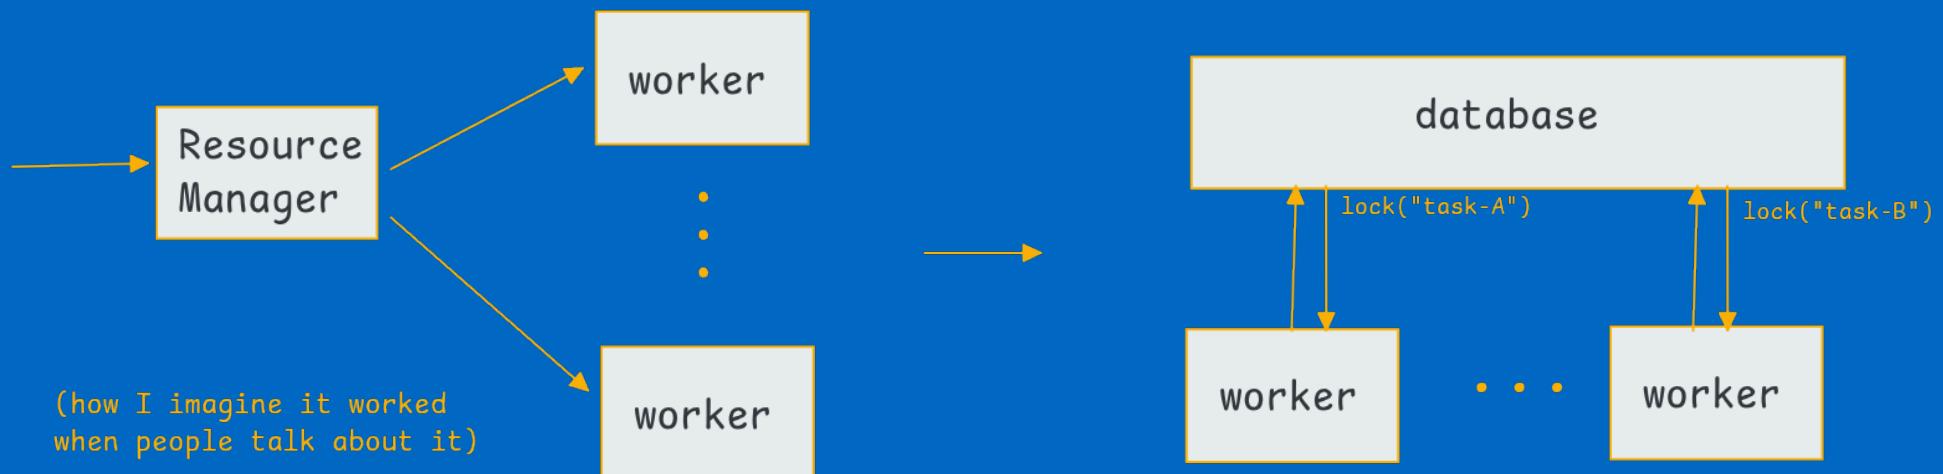
System in 2021

1. Distributed task pulling



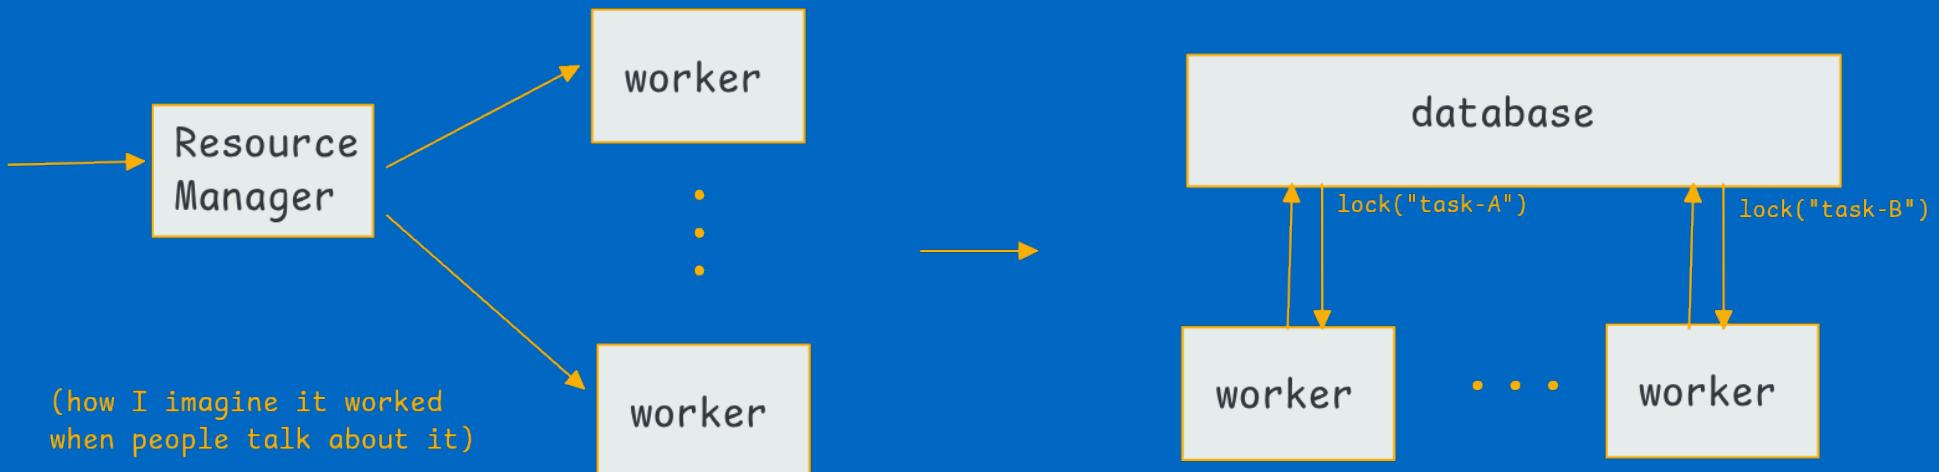
System in 2021

1. Distributed task pulling
 2. Single source of truth for tasks (Postgres)

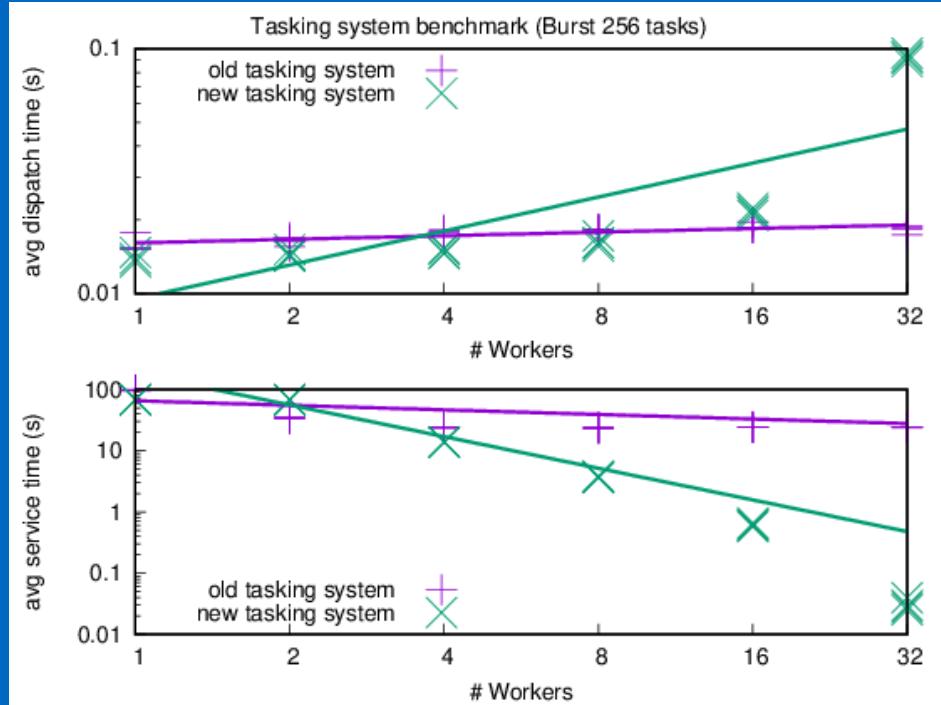


System in 2021

1. Distributed task pulling
 2. Single source of truth for tasks (Postgres)
 3. Auto resource management with PG Advisory Locks



System in 2021



- Better scalability
- More robust against race-conditions and deadlocks

Recap and Questions

Background • Pulp's Kitchen • New menu • Final words

Recap and Questions

- What Pulp is, it's architecture and operations that trigger tasks

Recap and Questions

- What Pulp is, it's architecture and operations that trigger tasks
- Some motivations for Pulp's tasking system

Recap and Questions

- What Pulp is, it's architecture and operations that trigger tasks
- Some motivations for Pulp's tasking system
- Some improvements over past tasking systems

Questions? Thoughts?

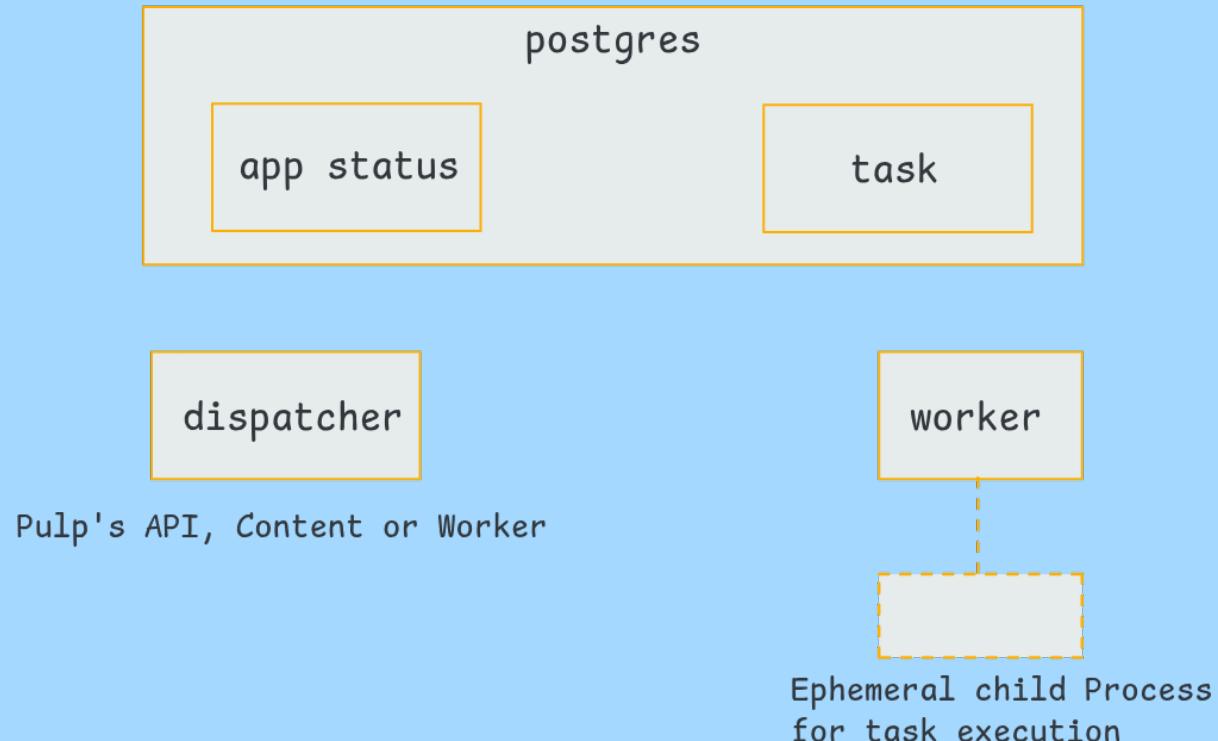
Reference

- Pulp's architecture:
 - <https://pulpproject.org/pulpcore/docs/admin/learn/architecture/>
- From celery to RQ
 - <https://pulpproject.org/blog/2021/06/21/a-new-tale---pulp-gets-queueless-distributed-task-workers/>
- From RQ to vanilla
 - <https://pulpproject.org/blog/2018/05/08/pulp-3-moving-to-rq-for-tasking/>
 - <https://www.youtube.com/watch?v=YWKw4RYluPM&pp=0gcJCSIKAYcqlYzv>

2. Cooking in Pulp's kitchen

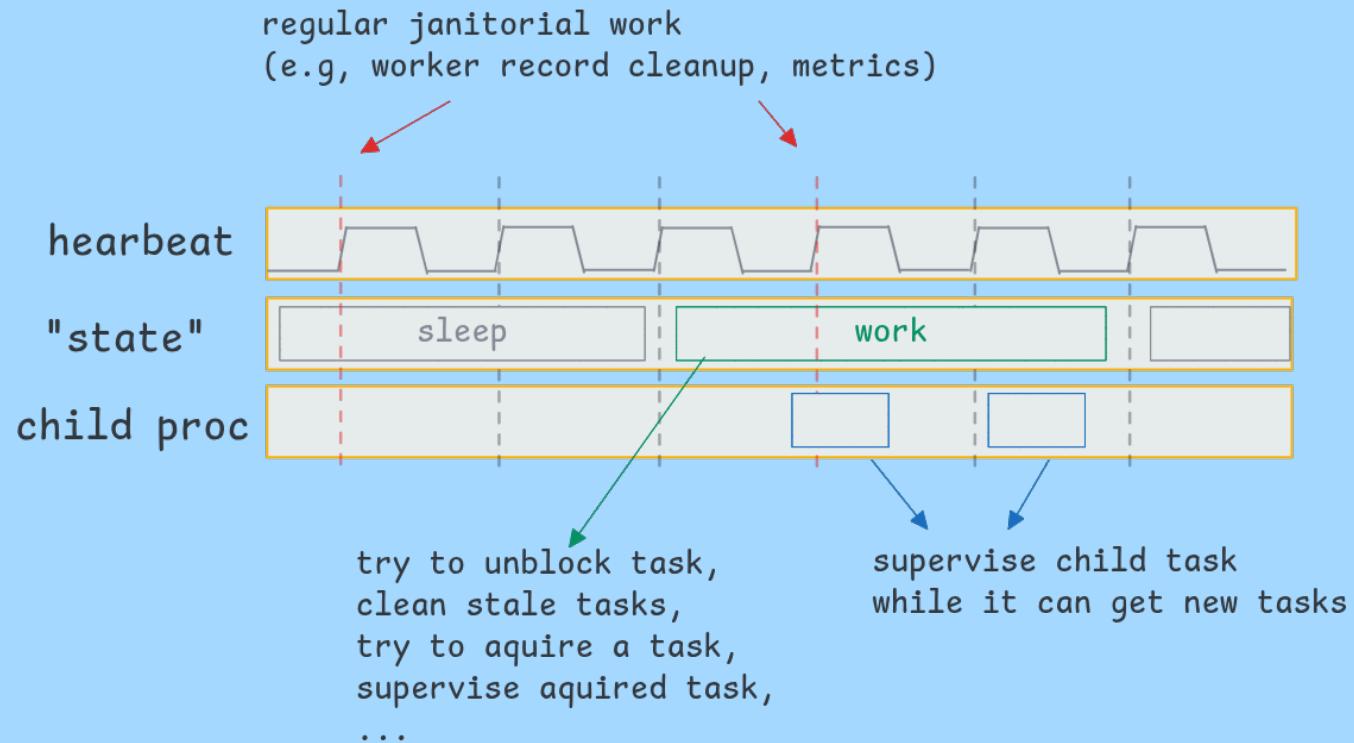
Background • Pulp's Kitchen • New menu • Final words

Components



Background • Pulp's Kitchen • New menu • Final words

Workers



Tasks

```
name = models.TextField()  
state = models.TextField(choices=TASK_CHOICES)  
reserved_resources_record = ArrayField(models.TextField(), null=True)  
unblocked_at = models.DateTimeField(null=True)  
immediate = models.BooleanField(default=False, null=True)
```

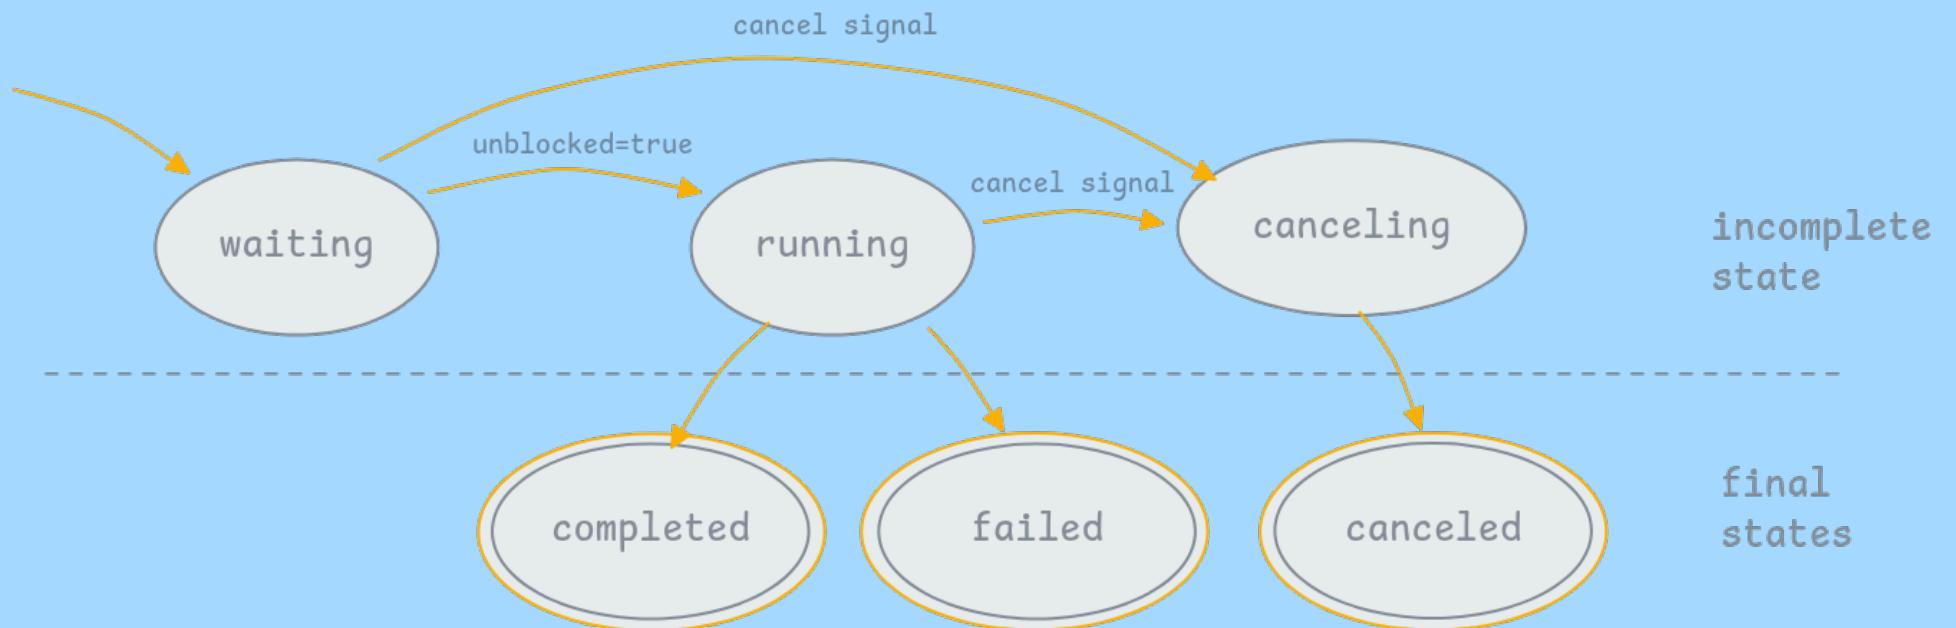
Tasks

```
name = models.TextField()  
state = models.TextField(choices=TASK_CHOICES)  
reserved_resources_record = ArrayField(models.TextField(), null=True)  
unblocked_at = models.DateTimeField(null=True)  
immediate = models.BooleanField(default=False, null=True)
```

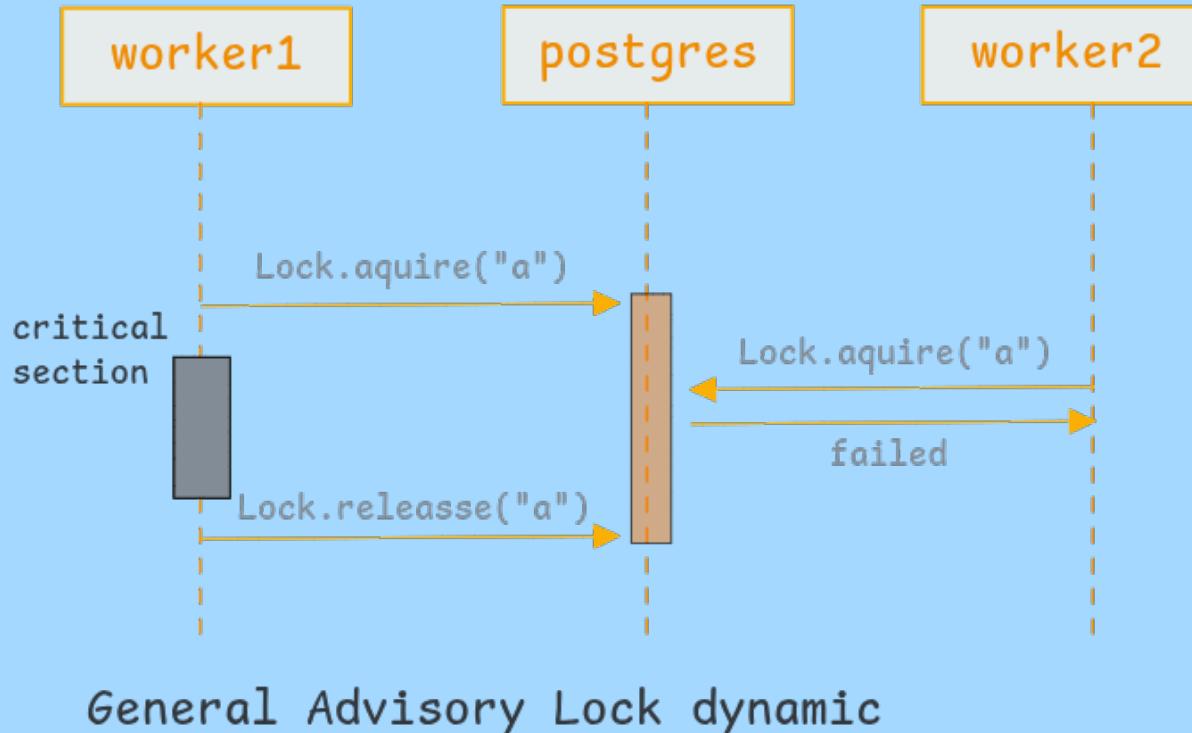
Tasks

```
name = models.TextField()  
state = models.TextField(choices=TASK_CHOICES)  
reserved_resources_record = ArrayField(models.TextField(), null=True)  
unblocked_at = models.DateTimeField(null=True)  
immediate = models.BooleanField(default=False, null=True)
```

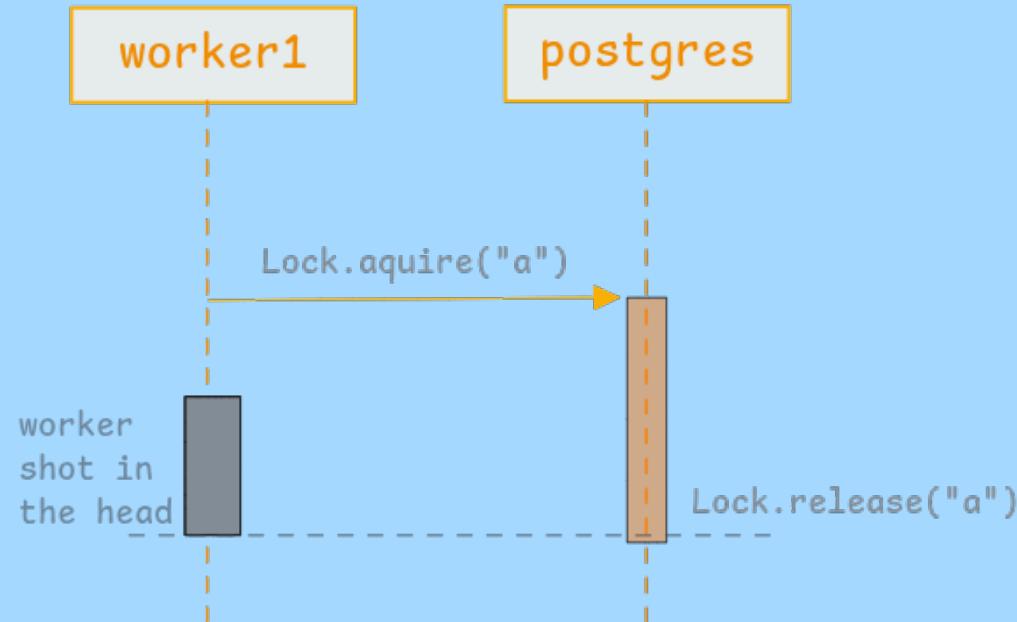
Tasks



Locks



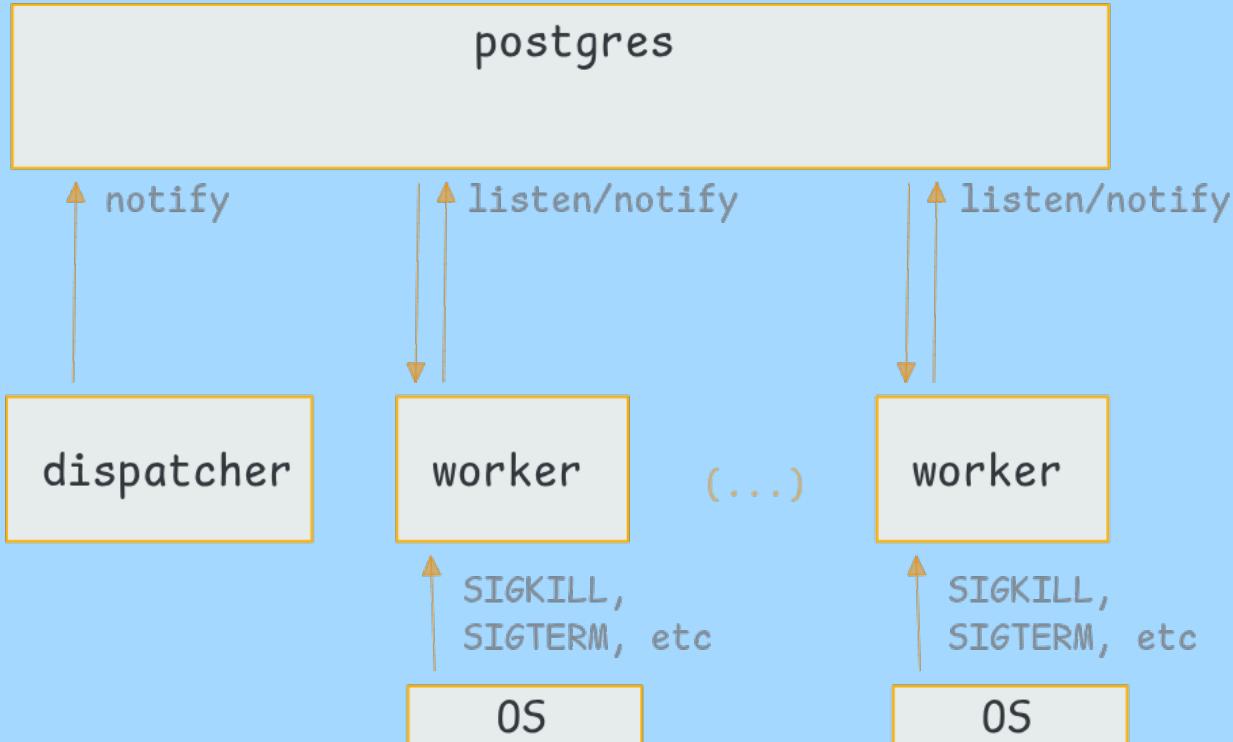
Locks



Session-based Advisory Locks

Background • Pulp's Kitchen • New menu • Final words

Signals



Signals

```
def _pg_notify_handler(self, notification):
    if notification.channel == "pulp_worker_wakeup":
        self.wakeup = True
    elif notification.channel == "pulp_worker_metrics_heartbeat":
        self.last_metric_heartbeat = datetime.fromisoformat(notification.payload)
    elif self.task and notification.channel == "pulp_worker_cancel":
        if notification.payload == str(self.task.pk):
            self.cancel_task = True
```

Recap and Questions

Background • Pulp's Kitchen • New menu • Final words

Recap and Questions

- Overview of tasking system components

Recap and Questions

- Overview of tasking system components
- Workers, tasks and locks

Recap and Questions

- Overview of tasking system components
- Workers, tasks and locks
- Inter-app communication with signals

Recap and Questions

- Overview of tasking system components
- Workers, tasks and locks
- Inter-app communication with signals

Questions? Thoughts?

Reference

- Documentation:
 - <https://pulpproject.org/pulpcore/docs/dev/learn/plugin-concepts/#tasks>
- Pulpcore's codebase:
 - pulpcore.app.models.task::Task
 - pulpcore.app.models.status
 - pulpcore.tasking.entrypoint
 - pulpcore.tasking.worker
 - pulpcore.tasking.tasking (specially dispatch)
- Obs: the contents of this section are roughly referring to pulpcore<3.85

3. New Menu

Background • Pulp's Kitchen • New menu • Final words

Requirements

- Pulp Services demand rapidly increasing. They need:
 - Throughput of 75-250k task/day (baseline 10k)
 - Connection pooling (Amazon's RDS Proxy)

Requirements

pulp-worker creates some database pressure when scaling beyond 250 workers #6616

[Edit](#) [New issue](#) [🔗](#)

[Open](#)

decko opened on May 26 Member ...

Version
3.78.0

Describe the bug
We ran a test to check the number of **idle** pulp-workers that a system could handle. We tested it with 50, 100, 250, 500 and 1000 workers.
The database instance is a AWS RDS Postgres, running on a db.m7g.2xlarge machine, with 8 Cores and 32GiB of RAM.
Each pulp-worker have access to one entire core, with 512MiB of RAM, which is not used most of the time.
The telemetry was disabled for this test(`OTEL_ENABLED=False`).
Each run consisted in scale the number of workers in a certain pace until it reaches the desired number and stay stable for 1 hour.

Assignees [⚙️](#)
No one - [Assign yourself](#)

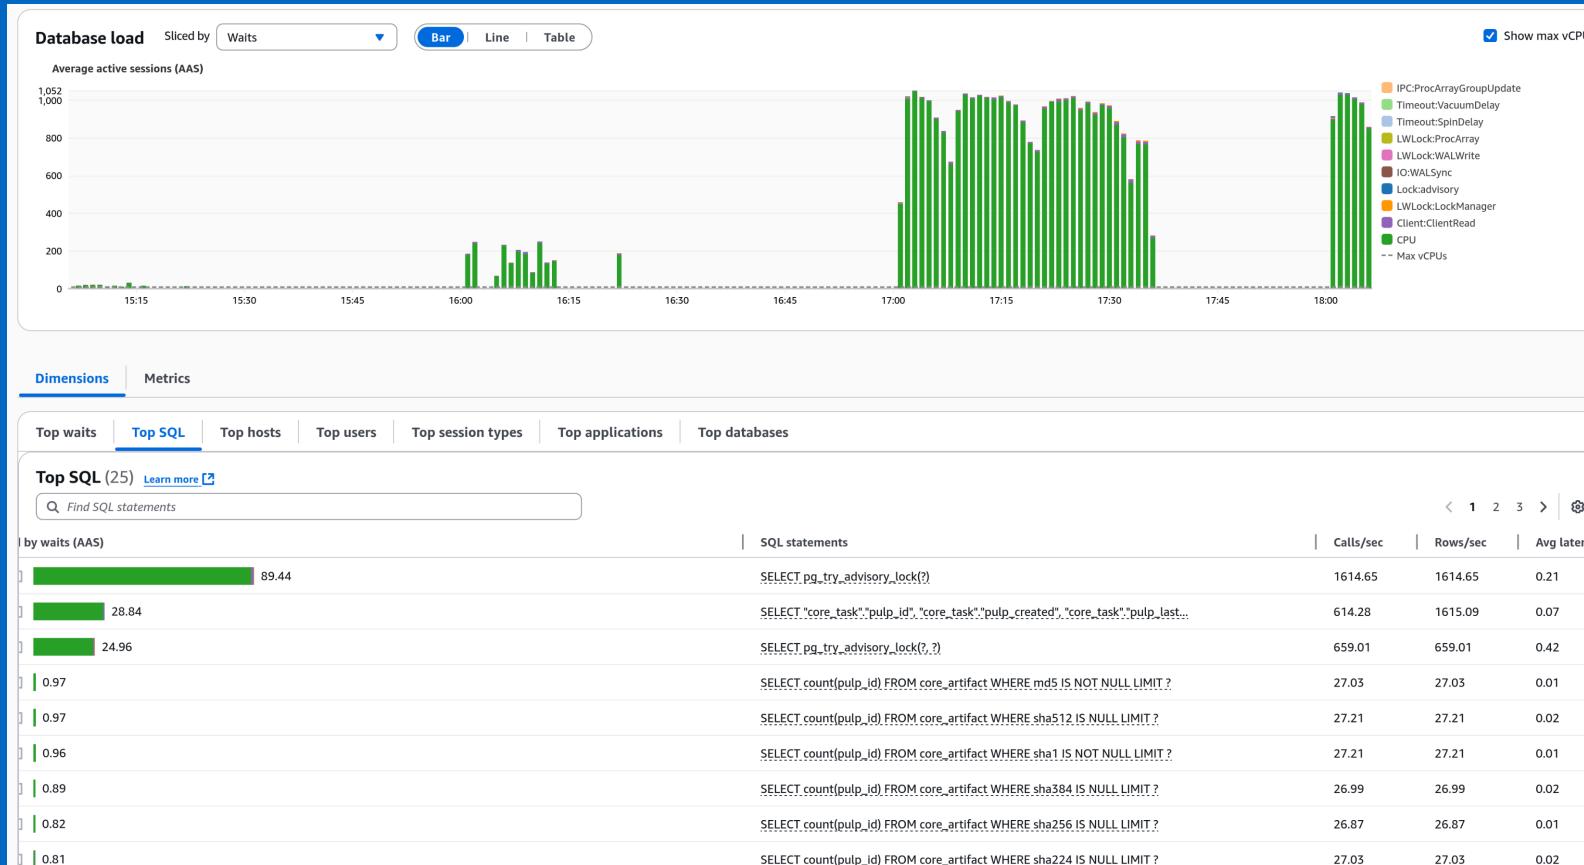
Labels [⚙️](#)
Issue

Type [⚙️](#)
No type

Projects [⚙️](#)
No projects

Milestone [⚙️](#)

Requirements



Requirements

Conditions that cause pinning for RDS for PostgreSQL

For PostgreSQL, the following interactions also cause pinning:

- Using SET commands.
- Using PREPARE, DISCARD, DEALLOCATE, or EXECUTE commands to manage prepared statements.
- Creating temporary sequences, tables, or views.
- Declaring cursors.
- Discarding the session state.
- Listening on a notification channel.
- Loading a library module such as auto_explain.
- Manipulating sequences using functions such as nextval and setval.
- Interacting with locks using functions such as pg_advisory_lock and pg_try_advisory_lock.

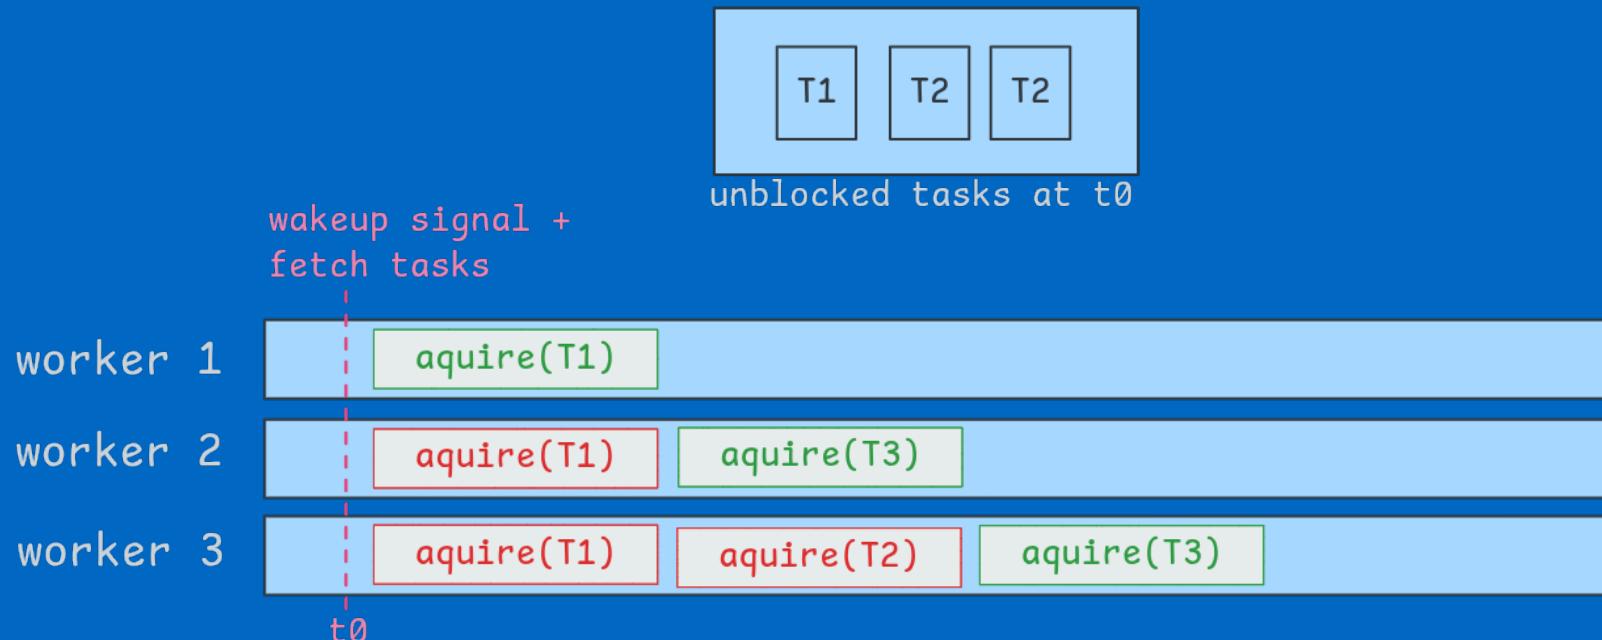
Note

RDS Proxy does not pin on transaction level advisory locks, specifically pg_advisory_xact_lock, pg_advisory_xact_lock_shared, pg_try_advisory_xact_lock, and pg_try_advisory_xact_lock_shared.

- Setting a parameter, or resetting a parameter to its default. Specifically, using SET and set_config commands to assign default values to session variables.
- Calling stored procedures and stored functions doesn't cause pinning. RDS Proxy doesn't detect any session state changes resulting from such calls. Make sure that your application doesn't change session state inside stored routines if you rely on that session state to persist across transactions. For example, RDS Proxy isn't currently compatible with a stored procedure that creates a temporary table that persists across all transactions.
- Discarding session state. If you use connection pooling libraries with DISCARD ALL query configured as a reset query, RDS Proxy pins your client connection on release. This reduces the proxy's multiplexing efficiency and might lead to unexpected results because the DISCARD ALL command can interfere with session state management.

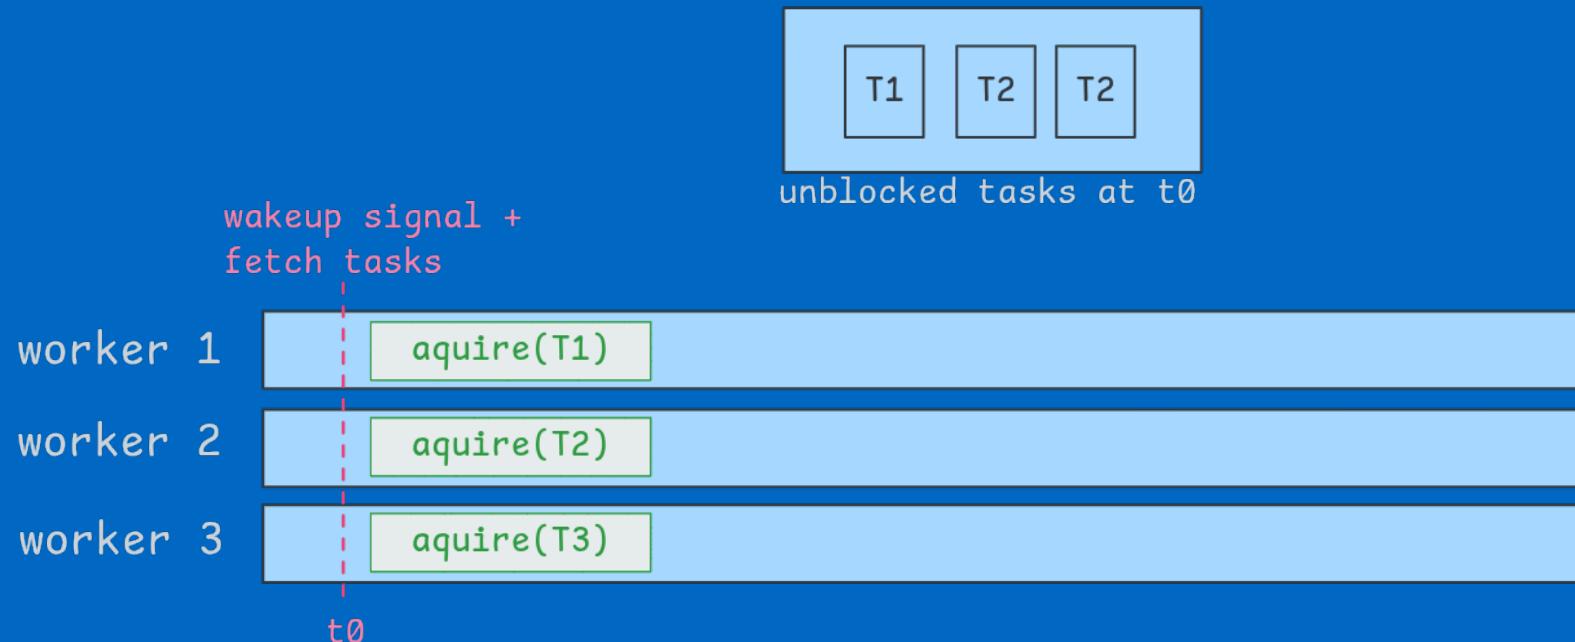
Changes: Worker's overhead

- Randomization on acquiring task



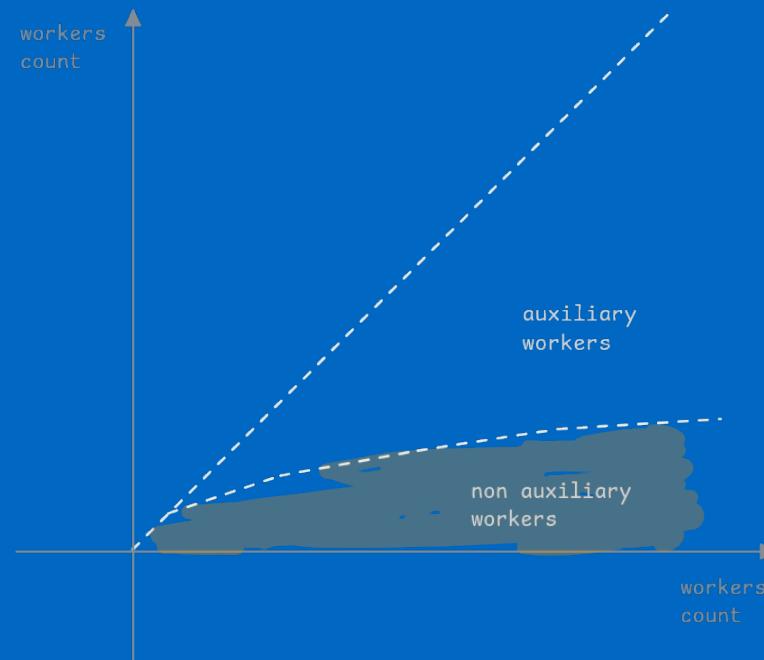
Changes: Worker's overhead

- Randomization on acquiring task



Changes: Worker's overhead

- Auxiliary workers: only supervise tasks



Changes: Worker's overhead

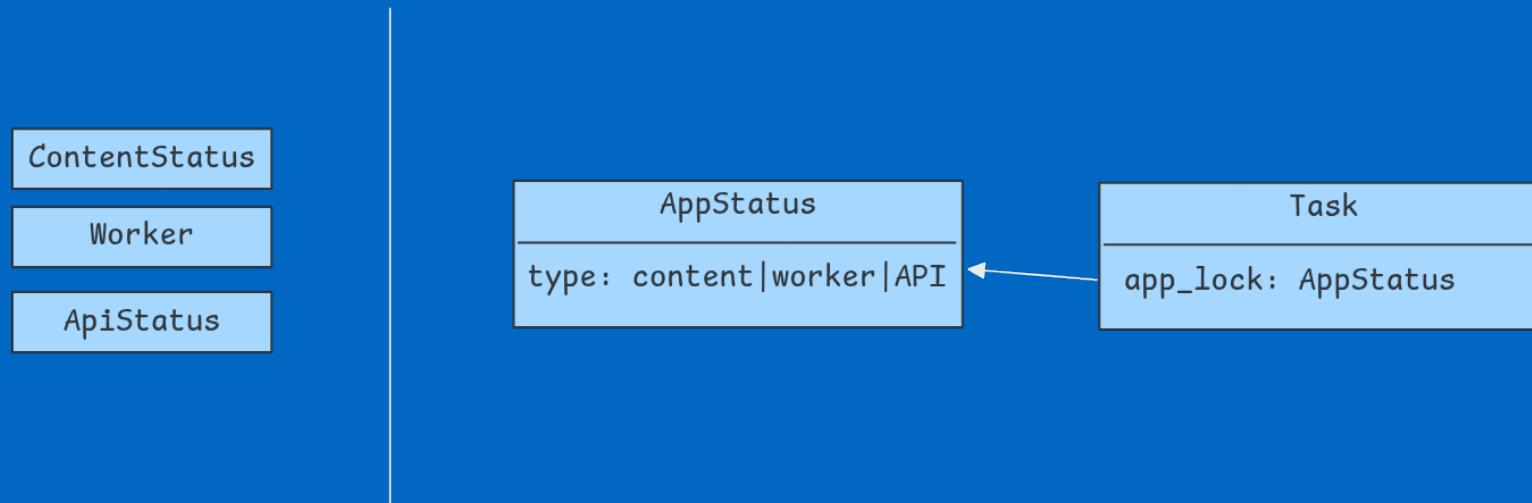
- Wakeup unblock/handle signals

dispatch → new waiting task → can be unblocked
can't be acquired

unblock → new unblocked task → can't be unblocked
can be acquired

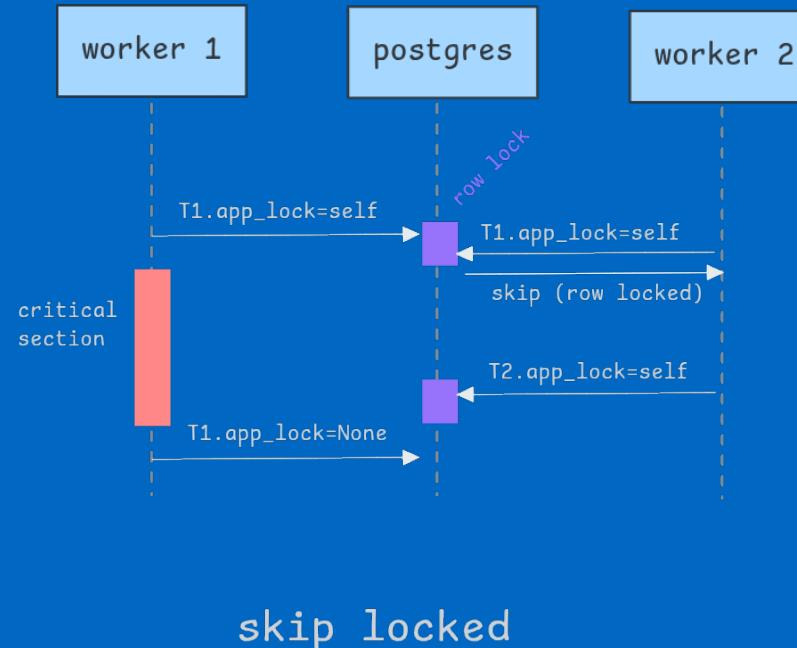
Changes: RDS Proxy support

- App Status consolidation
 - Required to build app_lock foreign key to any pulp app



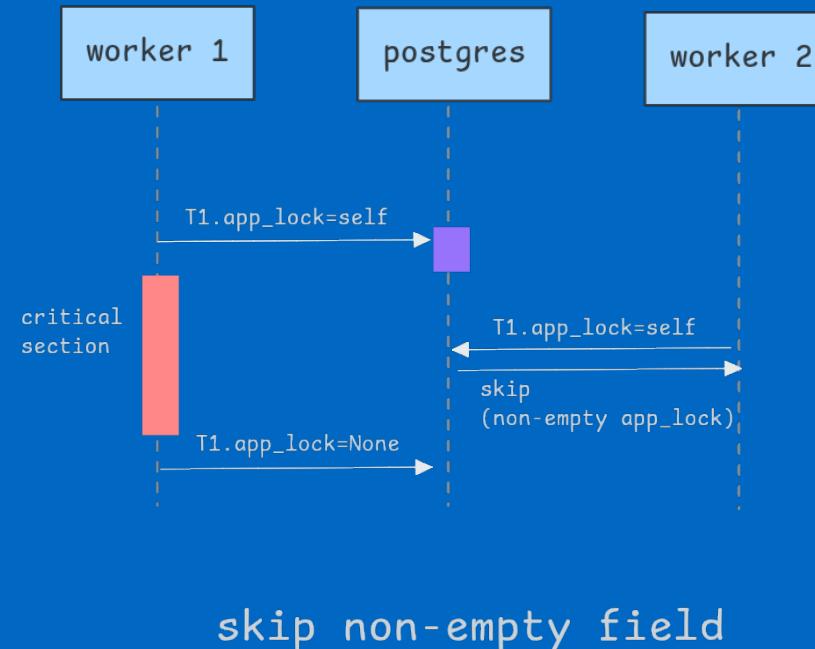
Changes: RDS Proxy support

- Custom PG update query to replace advisory locks



Changes: RDS Proxy support

- Remove session-based advisory locks



Recap and Questions

Background • Pulp's Kitchen • New menu • Final words

Recap and Questions

- Requirements from Pulp Services installation

Recap and Questions

- Requirements from Pulp Services installation
- Changes related to worker's overhead

Recap and Questions

- Requirements from Pulp Services installation
- Changes related to worker's overhead
- Changes related to providing support for connection pooling

Recap and Questions

- Requirements from Pulp Services installation
- Changes related to worker's overhead
- Changes related to providing support for connection pooling

Questions? Thoughts?

Reference

- Pulp Services' Jira
 - <https://issues.redhat.com/browse/PULP-674>
- Decko's issue with important results
 - <https://github.com/pulp/pulpcore/issues/661>
- See workers code for pulpcore>3.85

4. Final words

Background • Pulp's Kitchen • New menu • Final words

Conclusion

Background • Pulp's Kitchen • New menu • Final words

Conclusion

- This is *some* of the actual work done so far
 - ▶ ommited a lot of discussions, ideas, etc to fit time

Conclusion

- This is *some* of the actual work done so far
 - ▶ ommited a lot of discussions, ideas, etc to fit time
- Personal thought
 - ▶ lack of upstream instrumentation for understanding key metrics of the system

Conclusion

- This is *some* of the actual work done so far
 - ▶ ommited a lot of discussions, ideas, etc to fit time
- Personal thought
 - ▶ lack of upstream instrumentation for understanding key metrics of the system
- Kudos again to the tasking working group!

Conclusion

- This is *some* of the actual work done so far
 - ▶ ommited a lot of discussions, ideas, etc to fit time
- Personal thought
 - ▶ lack of upstream instrumentation for understanding key metrics of the system
- Kudos again to the tasking working group!

THE END